

# Edna: Data Disguises for Web Applications

Anonymous Authors

## 1 Setting and Goals

Our design considers typical database-backed applications that link a library implementing our disguising tool (Edna). Edna provides an API that the application uses to specify disguises, apply them, and (in some cases), reveal disguised data with proper authorization (Table 2). Edna executes the database transformations required as part of a disguise, and stores disguised data in a secure way, subject to the threat model we describe in §1.3.

### 1.1 Use Cases Edna Should Support

**Example: HotCRP.** We imagine three useful disguises that aid privacy in HotCRP:

1. Per-user GDPR Account Deletion, which removes all of a particular user’s data such as reviews and papers.
2. Universal Comment Removal, which removes all comments on papers.
3. Universal Decorrelation, which decorrelates users from their authored reviews and papers. Decorrelation of a paper or review (1) generates an anonymous user in the user table, and (2) rewrites the foreign key from the paper or review to the users table to point to the newly created anonymous user.

Users who invoke GDPR deletion are given the option to return by revealing their disguise; however, between disguising and revealing their account, no other entity learns the identity or existence of the user (ensuring that the disguise complies with the GDPR).

After several years, the admin can decay identifiers in the conference data by decorrelating all users from their authored reviews and papers. This helps protect reviewers and paper authors from being unblinded in the case of a data breach.

Even after universal decorrelation, it would be nice if HotCRP allowed users to apply GDPR account deletion to delete (even decorrelated) papers and reviews that they wrote:

a disguise that applies to a user’s data may, in some circumstances, need to apply to data decorrelated from the user.

Application of universal decorrelation also should not prevent the application of universal comment removal: the admin may decide, after applying universal decorrelation, that arbitrary text in paper comments also identifies their authors (which are now anonymous users), and invoke universal comment removal. This demonstrates how the developer can simply apply a new disguise to fix mistakes in a previously applied disguise that left some identifying information around, even if the new disguise now applies to data of anonymous users.

Finally, HotCRP would retain much of its usefulness if users could be permitted to view reviews on their papers, or edit their reviews, even if these papers and reviews have been decorrelated and belong to anonymous users. However, this should be possible *without* revealing to other users of the system which actual user authored these anonymized papers and reviews.

**Usability.** These disguising properties should be achieved with minimal burden on users.

When Edna disguises a user’s data for universal decorrelation or comment removal, Edna emails a unique URL (a *capability*) to the user corresponding to that a paper or review only known to the user. When the user wants to ask for permission to view (or edit, if permitted) the undisguised paper or review, the user simply has to click the URL link and log into the application. Thus, application of universal disguises should not require a user to be online in order to gain the capability to apply further disguises or view their disguised data.

Similarly, when Edna disguises a user’s data for GDPR deletion, the invoking user receives a URL that allows the user to reveal the disguise, restoring their account.

**General Use Cases.** From the HotCRP example, we generalize a handful of the most beneficial and practical use cases that Edna should support:

- **1st-Party GDPR-Compliant Disguising and Revealing.** Edna should support user-invoked disguise(s) to modify, decorrelate and/or delete the user’s data and account that enable applications to meet the requirements of the GDPR’s right to be forgotten. These disguises can optionally be revealable, which enables users to restore the original state of disguised data to e.g., permanently restore their accounts and/or data ownership.
- **3rd-Party Disguising.** Edna should support disguise applied by users (such as an admin) that modify, decorrelate, and/or delete data of other users in the system. For example, universal conference anonymization or comment removal alters data not associated with the invoking admin.
- **Temporary Recorrelation w/out Database Changes.**
  - *For Disguising Decorrelated Data:* A disguise targeting a user’s data can, if authorized, also operate on data decorrelated from the user by prior disguises without modifying the database’s persistent links between data and owning users. For example, GDPR account deletion removes user  $u$ ’s reviews even if Edna has decorrelated these reviews from  $u$ .
  - *For Per-User Revealed Views and Permissions:* Users should be able to operate with ownership permissions on data previously correlated with their account, but decorrelated by a disguise. Users do so by interacting with a personalized, temporary database view when they visit their decorrelated data. This should require no modifications to the actual database contents.  
For example, a user whose authored papers have been decorrelated can log in and view their papers and reviews on their papers.
- **Disguising Anonymized Users.** Edna can disguise data of anonymized users. For example, universal comment removal may remove the data of anonymized users generated by universal decorrelation. This allows disguises to be developed over time, and remove identifying data potentially mistakenly missed by prior disguises, but which now belongs to anonymous users.

## 1.2 Edna Abstractions

Edna supports these use cases by exposing three abstractions to applications:

- Application *principals*, which correspond to users of the application, and which are uniquely identifiable (e.g., via a user ID). We denote an application principal as  $p \in P$ , where  $P$  is the set of all application principals.

- Application *pseudoprincipals*, which correspond to anonymous users not tied to any natural person, and which are uniquely identifiable (e.g., via a user ID). Pseudoprincipals are part of the set of application principals  $p \in P$ . A pseudoprincipal may or may not be indistinguishable from a real principal: for example, it may not have an associated password.
- A disguise  $d \in D$  invoked by some principal  $p$ , which transforms the application database.

Edna operates on two categories of data:

1. **Database Contents:** The application reads and writes database contents, which include visible, undisguised and currently-disguised data. Edna modifies the data database contents when it applies and reveals disguises.
2. **Disguise Diffs:** Edna records the set of disguise diffs—which contain the original database contents and the modifications performed to them—and uses diffs to compose disguises on top of one another, or to reveal the updates done by a disguise. Every disguise diff has an associated principal (namely the one associated with the modified data).

**Data Access Control.** The application’s normal permissions logic controls data access to database contents. The application invokes Edna to access and needs to parse diff contents directly: Edna exposes an API (§2) allowing the application to learn specific semantic information contained in diffs (e.g., whether a particular principal used to own a now-decorrelated piece of data) if provided appropriate authorization. This authorization also allows the application to ask Edna to use diff information to reveal or apply disguises.

Two types of *capabilities* control access to diffs:

1. **Data Capability**  $C_p^{\text{data}}$ : Grants read access to disguise diffs associated with principal  $p$ .
2. **Locating Capability**  $C_{pd}^{\text{loc}}$ : Allows locating the disguise diffs associated with principal  $p$  produced from applying  $d$ , but grants no access to the diffs’ data.

Edna needs both the data and locating capabilities ( $C_p^{\text{data}}, C_{pd}^{\text{loc}}$ ) to access disguise diffs of disguise  $d$ . Each capability alone is insufficient: an attacker who can spoof  $C_{pd}^{\text{loc}}$  will not be able to access disguise diffs without  $C_p^{\text{data}}$ , although they can learn that  $d$  disguised some of  $p$ ’s data; and an attacker holding  $C_p^{\text{data}}$  cannot locate the disguise diffs to access without  $C_{pd}^{\text{loc}}$ .

## 1.3 Threat Model

Data disguises protect user information in a web application against external observation and service compromise. An

external observer is a user of the web application (authenticated or unauthenticated) who observes information exposed through using the application. A service compromise occurs when an attacker compromises the web application and gains full access to the server. The attacker therefore can access any data stored, perform any actions the application can perform, and access any information available to Edna.

A data disguise guarantees that the disguised data is hidden from any future attackers unless explicitly revealed by an authorized principal. In particular, an attacker who compromises Edna at time  $t$  learns *nothing but*:

1. the (plaintext) contents of the application database at or after time  $t$ ;
2. the disguises invoked, and the identity of the principals invoking them, after time  $t$ ; and
3. all of a principal’s disguise diffs acquired prior to time  $t$ , should the principal authorize access to disguise diffs for purposes of revealing or performing some application action after time  $t$ .

We make standard assumptions about the security of cryptographic primitives: attackers cannot break encryption and keys stored with non-colluding clients are safe.

Edna operates in an honest-but-curious setting: even if compromised, Edna faithfully executes its protocols, but exposes all data accessed to the attacker.

## 1.4 Security Goals

With this threat model, Edna seeks to meet four security goals:

**(1) Authorized Disguises.** Only a client properly authenticated as a principal  $p$  who is authorized to invoke  $d$  can apply (and later reveal)  $d$ .

**(2) Secure Disguise Diffs.** Only an authenticated principal who provides both locating capability  $C_{pd}^{loc}$  and data capability  $C_p^{data}$  can authorize access to the disguise diffs, as well as any information (e.g., ownership permissions) derived from diffs.

**(3) Privacy of Disguise History.** An attacker cannot learn the set of disguises that have disguised a principal’s data. An attacker only learns a disguise  $d$  exists if an authenticated principal  $p$  provides locating capability  $C_{pd}^{loc}$  to Edna.

**Non-Goals.** Under Edna’s threat model, the following properties are out of scope:

- Security of disguised data when an attacker has prior snapshots of the system.
- Security of metadata such as the number of disguises applied to, or number of disguise diffs associated with, *some* principal (where the principal ID is unknown).
- Security of metadata when an attacker has  $C_{pd}^{loc}$ : the attacker can learn that  $d$  applied to a particular principal  $p$ , and the number of disguise diffs for that  $p$  and  $d$ .

- Privacy guarantees about undisguised data: if identifying data about  $p$  is not covered by disguise specification  $d$ , it remains visible to attackers after corresponding instance  $d$  applies. For example, a disguise that leaves contents of posts unmodified does not hide identifying references to users in that content.
- Hiding whether data is disguised or not. For example, if a post’s author is anonymized as “anonFox”, Edna leaks the fact that a post’s author has been disguised.

## 2 Edna API.

Edna’s API in Table 2 enables the application to support disguise application and revealing. The application also uses Edna to access and interpret diffs authorized by the provided capabilities. Applications integrating with Edna may modify the client API so that the client can pass capabilities to the application along with the action to perform. For example, the application can send users URLs to click that invoke application actions with capability arguments.

We next describe example usages of the API.

**Principal Registration.** Every client who creates an application account for principal  $p$  registers an email and public key with Edna. Edna remembers each public key  $k_p^{pub}$  along with  $p$ ’s ID.

**Disguising and Revealing.** When a client speaking for principal  $p$  wants to apply a disguise  $d$ , the application invokes `ApplyDisguise` with the client-provided data capability and a set of locating capabilities. Each locating capability  $C_{pd'}^{loc}$  allows Edna to access  $p$ ’s disguise diffs from disguise  $d'$ , and to compose the new disguise on top of these disguises during application. Applying a disguise returns the locating capability  $C_{pd}^{loc}$  that the application sends to the client (via e.g., emailing the client a link).

In order to reveal data disguised by  $d$  for principal  $p$ , the application invokes `RevealDisguise` on behalf of a client speaking for  $p$ . The client provides its data capability  $C_p^{data}$  and locating capability  $C_{pd}^{loc}$ . The application passes these to Edna to restore the original, undisguised data with access to  $p$ ’s diffs corresponding to  $d$ .

**Performing Permitted Application Actions.** Imagine that HotCRP disguise  $d$  anonymizes all reviewers and authors of a conference by decorrelating them from their papers and reviews respectively. As described above, HotCRP emails the client a link when that their data has been decorrelated, which contains the locating capability in the URL. Assume that HotCRP sends a unique link for every piece of data (e.g., every review or paper); HotCRP could also send a single link for all decorrelated data.

The application normally ensures that only a client speaking for  $p$  should have read-only access to reviews for  $p$ ’s authored papers, and read-write access to reviews  $p$  wrote.

Symbol	Description
$p$	application principals, corresponding to a user ID in the application
$d$	disguise invoked by authorized principals that transforms the application database in privacy-preserving ways
$\text{op}_d$	a disguise operation (either a removal, modification, or decorrelation)
$\Delta_{pd}$	database diff associated with $p$ produced by disguise $d$ .
$C_p^{\text{data}}$	data capability that grants access to all disguises' diffs associated with principal $p$ .
$C_{pd}^{\text{loc}}$	locating capability that allows the holder to find the database diffs from $d$ associated with a principal $p$ .
$k_p^{\text{pub}}$	public key of $p$
$k_p^{\text{pri}}$	private key of $p$

Table 1: Notation used to describe Edna's design.

API Call	Description
<code>RegisterPrincipal(<math>p</math>, Email email, PubKey <math>k_p^{\text{pub}}</math>) <math>\rightarrow ()</math></code>	Registers $p$ as a principal with public key $k_p^{\text{pub}}$ .
<code>ApplyDisguise(<math>p</math>, Disguise <math>d</math>, DataCap <math>C_p^{\text{data}}</math>, Vec&lt;LocCap&gt; lcaps) <math>\rightarrow \text{Vec&lt;LocCap&gt;}</math></code>	Applies disguise $d$ and returns the corresponding disguise $d$ and locating capabilities $C_p^{\text{loc}}d$ . The argument capabilities allow Edna to locate and access disguise diffs, which Edna uses to compose disguise $d$ on top of prior disguises.
<code>RevealDisguise(<math>p</math>, Disguise <math>d</math>, DataCap <math>C_p^{\text{data}}</math>, LocCap <math>C_{pd}^{\text{loc}}</math>) <math>\rightarrow ()</math></code>	Reveals disguise $d$ . The argument capabilities grant Edna access to database diffs from applying $d$ , which allows Edna to reveal undisguised data.
<code>CapEstablishesOwnership(<math>p</math>, DatabaseObj <math>x</math>, DataCap <math>C_p^{\text{data}}</math>, LocCap <math>C_{pd}^{\text{loc}}</math>) <math>\rightarrow \text{bool}</math></code>	Returns whether principal $p$ has ownership rights to the provided database object. The capabilities allow Edna to access diffs that can prove that $p$ had been correlated with the object.

Table 2: The Edna Library API

If the client speaking for  $p$  wants to edit a review as it would normally, the application (post- $d$ ) will not grant the client any read/write permissions to perform actions on any reviews of the conference: any previously correlated paper or review is now associated with pseudoprincipal  $q$ .

Instead, for the client speaking for  $p$  to regain authorship or reviewership permissions of a particular review or paper:

- The client clicks on the URL emailed to the client for that piece of data when it was decorrelated, which opens a HotCRP page.
- Javascript on the webpage retrieves the client's data capability stored by the client's browser, and extracts the locating capability and the data object the client wants to access from the URL.
- The client then must log into HotCRP as principal  $p$ .
- HotCRP invokes Edna's `CapEstablishesOwnership` for the data object in question, using the capabilities

passed in by the client.

- If Edna returns `true`, the application permits the client to view the object with ownership rights (e.g., ability to see reviews on authored papers, or edit their review); otherwise, the application does not allow the client to perform privileged actions.

## 3 Design

### 3.1 Disguise Operations.

Because disguises are inherently application-specific, the application developer specifies a disguise as a set of predicated disguise operations  $\text{op}_d$  to perform.

Operations  $\text{op}_d$  of disguise  $d$  take data objects as input and execute updates to application data. Edna automatically generates database change records when applying  $\text{op}_d$ .

Developers describe which principal(s) an operation's generated database change record corresponds to. For example, a database change record generated by removing comment

may correspond to the principal whose ID is referenced by the author column.

Developers also specify an application-aware pseudoprincipal generation policy, namely how to generate new user accounts in a manner that the application can handle (e.g., pseudoprincipals may not have email addresses).

Finally, developers specify which principals are authorized to apply the disguise: enforcing access control for disguising is left to the application.

Operations come in three forms:

1. Modify: change an attribute of the data object.
2. Remove: delete the data object.
3. Decorrelate: generate a *pseudoprincipal*  $q$ , and rewrite the foreign key to original principal  $p$  from the data object to instead point to  $q$ .

For each  $\text{op}_d$ , the application developer specifies:

- An associated predicate over the application database that selects  $\text{op}_d$ 's input objects in a SQL-like fashion.
- The type of operation and its arguments (e.g., which attributes of the data object to modify).
- The corresponding principal(s) that should have the capability to access  $\text{op}_d$ 's generated disguise change or correlation record.

### 3.2 Database Diffs.

Every  $\text{op}_d$  produces a *database diff*  $\Delta_{pd}$  associated with the disguise  $\delta$  and a principal  $p$ . Each  $\Delta_{pd}$  contains the ID of the disguise, the associated principal  $p$ 's ID, and a random nonce. In addition to this data, a decorrelation  $\Delta_{pd}$  contains the created pseudoprincipal's ID; a removal  $\Delta_{pd}$  contains the removed object's value; and a modification  $\Delta_{pd}$  contains the old and new value of the modified object.

### 3.3 Diff Access Control.

**Securing Access to Database Diffs.** Edna's design uses principal  $p$ 's private key  $k_p^{\text{pri}}$  as the data capability  $C_p^{\text{data}}$ . Edna secures  $p$ 's diffs  $\Delta_{pd}$  by encrypting them with  $k_p^{\text{pub}}$ . The diff's nonce ensures safety against known-plaintext attacks. Only a client who knows  $k_p^{\text{pri}}$  can access diffs.

**Securing Disguise History.** Edna's design may leak information that  $d$  disguised principal  $p$  because an adversary may learn that ciphertexts for  $\Delta_{pd}$  exist for a particular  $p$  and  $d$ . This is a problem: for example, Edna should not store information that a principal  $p$  has invoked GDPR deletion.

To avoid this problem, Edna stores an array of  $\Delta_{pd}$  ciphertexts at random location pointed to by locating capability  $C_{pd}^{\text{loc}}$ . This dissociates one disguise applied to a principal  $p$  from all other disguises, so an adversary only ever learns about a

single disguise if Edna gains permission to reveal or compose upon that disguise.

Note that an adversary without access to any  $C_{pd}^{\text{loc}}$  can learn that  $n$  diffs exist for *some*  $p$  and  $d$ , but cannot directly identify which  $p$  or  $d$ . If an adversary has access to  $C_{pd}^{\text{loc}}$ , the adversary can learn that  $d$  applied to  $p$ , and the number of  $\Delta_{pd}$  diffs for that  $p$  and  $d$ . This metadata is out of scope of our threat model. [lyt: An alternative design might remove the encrypted data completely (and email it to the client), so that an adversary doesn't even know of its existence. However, since we're allowing disguised data to be distinguished from undisguised data, this seems potentially unnecessary.]

**Storage of Capabilities.** Edna should not store  $C_p^{\text{data}}$  or  $C_{pd}^{\text{loc}}$ : an adversary could then learn that  $d$  has applied to principal  $p$ . Thus, these capabilities must be stored externally, even when no client speaking for  $p$  is currently online.

Clients already hold  $C_p^{\text{data}}$ : a client generates a keypair for a new principal  $p$  and stores  $k_p^{\text{pri}}$ , while registering  $k_p^{\text{pub}}$  with Edna. Edna only stores  $k_p^{\text{pub}}$ .

Because Edna generates  $C_{pd}^{\text{loc}}$  during disguising, Edna must communicate  $C_{pd}^{\text{loc}}$  to clients. Edna emails  $C_{pd}^{\text{loc}}$  to the corresponding email address associated with  $p$ . This allows Edna to disguise data of principals even when no client authenticated as  $p$  has an active session open.

### 3.4 Current Design: Discussion

As the current design stands, Edna supports 1st-person and 3rd-person disguising, as well as 1st-person revealing. Application developers can write GDPR-compliant and/or universal disguises with the primitives exposed by Edna to write disguise specifications. Edna uses diffs produced from disguising to reveal data when authorized to do so, and when revealing does not revert updates made to the data since the time of disguise application.

Edna also supports "Temporary Recorrelation without Database Changes" because Edna can determine the original owner of data as long as Edna has access to decorrelation diffs (which the client or the application provides via the appropriate capabilities). Edna can use information from decorrelation diffs to disguise decorrelated data as if it were owned by the original user; and Edna can support the API discussed in §2, which allows applications to query Edna to check ownership properties and grant authorized users personalized views and permissions to access data objects.

Furthermore, Edna does this while meeting all security goals: Edna supports authorized disguises and ensures the security of ownership claims, disguise diffs, and disguise history.

Our current design falls short, however, by failing to support "Disguising Anonymized Users" which we explain next.



### 3.5 Disguising Anonymized Users

To apply one disguise on top of another in a way that supports, e.g., comment removal after universal decorrelation, Edna must support *pseudoprincipal diffs*, namely diffs associated with pseudoprincipals.

These diffs can occur after at least one disguise has been applied: subsequent disguises may associate diffs with pseudoprincipals. For example, after universal decorrelation, every paper or review correlates with a pseudoprincipal user. When Edna applies universal decorrelation, decorrelation of all papers and reviews generates correlation diffs for these pseudoprincipals.

However, we now have a problem. Diffs are encrypted with  $C_p^{\text{data}}$  ( $p$ 's private key) and stored at  $C_{pd}^{\text{loc}}$ , and Edna emails  $C_{pd}^{\text{loc}}$  to (potentially offline) clients and does not retain  $C_{pd}^{\text{loc}}$ . However, pseudoprincipals have no corresponding real user (and, for unlinkability, Edna cannot store which pseudoprincipal-user correspondences). Thus, Edna has no way to communicate  $C_{pd}^{\text{loc}}$  to a user who speaks for pseudoprincipal  $p$ !

We have three potential approaches to solve this problem. Two we see as strawmen because they fail to meet our security and use case goals respectively:

1. *Weak Security*: Edna can store  $C_{pd}^{\text{loc}}$  for pseudoprincipal  $p$ . This means an adversary will be able to learn  $p$ 's undisguised data, even if it had been disguised by  $d$ .
2. *Permanent Disguises*: Edna throws away  $C_{pd}^{\text{loc}}$  for pseudoprincipal  $p$ . This means that the disguise modifications are permanent, and no links (to determine ownership) between  $p$  and other pseudoprincipals can be made. This affects temporary recorrelation, because decorrelation links between principals cannot be derived from discarded diffs.

We next describe a third option, which meets our security and use case goals while trading off efficiency and usability.

**Pseudoprincipal Private Keys.** When a pseudoprincipal  $q$  is generated to decorrelate data from  $p$ :

1. Edna generates keypair  $(k_q^{\text{pub}}, k_q^{\text{pri}})$  for  $q$
2. encrypts the private key of  $q$  with  $k_p^{\text{pub}}$  to produce  $\text{Enc}(k_q^{\text{pri}}, k_p^{\text{pub}})$
3. stores  $\text{Enc}(k_q^{\text{pri}}, k_p^{\text{pub}})$  at location  $C_{pd}^{\text{loc}}$  along with  $p$ 's encrypted diffs from disguise  $d$

This ensures that a client with  $C_{pd}^{\text{loc}}$  and  $C_p^{\text{data}}$  ( $k_p^{\text{pri}}$ ) can access  $k_q^{\text{pri}}$ , the data capability of pseudoprincipal  $q$ . Furthermore, only a client who holds  $C_{pd}^{\text{loc}}$  can learn that  $p$  has decorrelated data because associated pseudoprincipal private keys are not stored publically with  $p$ 's ID.

When Edna applies disguise  $d'$  to further disguise pseudoprincipal  $q$ 's data, Edna:

1. stores disguise diffs  $\Delta_{qd'}$  at location  $C_{qd'}^{\text{loc}}$
2. encrypts  $C_{qd'}^{\text{loc}}$  with  $k_p^{\text{pub}}$  to produce  $\text{Enc}(k_p^{\text{pub}}, C_{qd'}^{\text{loc}})$
3. stores  $\text{Enc}(k_p^{\text{pub}}, C_{qd'}^{\text{loc}})$  mapped to by principal  $q$ 's ID

Note that Edna stores  $C_{qd'}^{\text{loc}}$  ciphertexts associated with  $q$ 's ID, whereas for a real principal  $p$ , Edna does not store any similar metadata indicating how many disguises have applied to  $p$ . This means that an adversary can learn that *some* disguise has applied to and decorrelated data from  $q$ . However, this falls out of scope in our threat model:  $q$  is a pseudoprincipal created from a decorrelation operation, and therefore any disguise metadata regarding  $q$  has already been dissociated from the original principal  $p$ 's identity.

**Extending capabilities to authorize pseudoprincipal diffs access.** With this extended design, a client speaking for  $p$  that provides data and locating capabilities ( $C_p^{\text{data}}, C_{pd}^{\text{loc}}$ ) authorizes access to *both*  $p$ 's diffs for disguise  $d$ , and diffs of a pseudoprincipal  $q$  created by decorrelating  $p$ 's data. [lyt: I don't think we're necessarily allowing clients to speak for  $q$ , but rather to perform actions using  $q$ 's diffs.]

Given a data capability  $k_p^{\text{pri}}$ , Edna can access  $k_q^{\text{pri}}$  by decrypting  $\text{Enc}(k_p^{\text{pub}}, k_q^{\text{pri}})$ . With  $k_q^{\text{pri}}$ , Edna can decrypt all  $\text{Enc}(k_q^{\text{pub}}, C_{qd}^{\text{loc}})$  so that Edna has both the data and location capabilities for all of  $q$ 's disguise diffs. This allows the application to reveal, disguise, or perform application actions with access to diffs from all disguises applying to  $q$ .

While clearly more expensive for Edna to execute, this protocol allows Edna to recursively disguise pseudoprincipals, yet still require only the top-level original principal capabilities to support the desired use cases. Furthermore, this design still meets our security goals: data capabilities secure all disguise diffs, and an adversary learns no disguise metadata for real users.

## 4 Disguising and Revealing Semantics

This section describes the semantics of disguising provided by Edna. In particular, we describe the end state of a data object given some history of disguise applications and reveals.

**Some Notation.** We describe disguise histories as a list of  $A(d_i)$  and  $R(d_i)$  actions, where  $A(d_i)$  corresponds to the application of disguise  $d_i$ , and  $R(d_i)$  corresponds to the reveal of  $d_i$ 's diffs. Time moves to the right in the list.

For every data object  $x$  in the system,  $x_{\text{start}}$  describes its initial state, and  $x_{[A(d_1), \dots]}$  describes its state after Edna has applied the history  $[A(d_1), \dots]$ .

**Composing Multiple Disguise Applications.** Let  $d_1$  and  $d_2$  be two disguises, where  $d_1$  is applied after  $d_2$ . to produce disguise history  $[A(d_1), A(d_2)]$ . Let  $x$  be some application data object, where  $x_{\text{start}}$  is its initial state, and  $x_{[A(d_1), A(d_2)]}$  is the final state after Edna applies  $d_1$  and  $d_2$ . If both  $d_1$  and  $d_2$

update  $x$ , what is  $x_{[\mathbf{A}(d_1), \mathbf{A}(d_2)]}$ ? We consider two possible end states:

( $\mathcal{S}_1$ )  $x_{[\mathbf{A}(d_1), \mathbf{A}(d_2)]}$  reflects the application of *both*  $d_1$  and  $d_2$  to  $x_{\text{start}}$ .

In other words, if an  $\text{op}_{d_2}$ 's predicate  $\text{pred}$  matches  $x_{\text{start}}$ , then  $\text{op}_{d_2}$  is applied to  $x_{[\mathbf{A}(d_1)]}$ , even if  $\text{pred}$  does not match  $x_{[\mathbf{A}(d_1)]}$ .

$\text{op}_{d_2}$  updates are applied sequentially after  $d_1$ 's updates if the two modify the same object attribute.

( $\mathcal{S}_2$ )  $x_{[\mathbf{A}(d_1), \mathbf{A}(d_2)]}$  reflects the application of  $d_1$  to  $x_{\text{start}}$ , followed by the application of  $d_2$  to  $x_{[\mathbf{A}(d_1)]}$ . When deciding whether to apply to  $x$ , Edna matches  $\text{op}_{d_2}$ 's predicate only against  $x_{[\mathbf{A}(d_1)]}$ , and not against the original state  $x_{\text{start}}$ .

For example, let  $\text{op}_{d_1}$  decorrelate all posts from authors predicated on  $\text{author} = \text{Bea}$ , and let  $\text{op}_{d_2}$  remove all posts predicated on  $\text{author} = \text{Bea}$ .

( $\mathcal{S}_1$ ) Both  $\text{op}_{d_1}$  and  $\text{op}_{d_2}$  update posts originally having author "Bea", resulting in the removal of all posts originally with author Bea.

( $\mathcal{S}_2$ ) Applying  $\text{op}_{d_1}$  results in all posts with author "Bea" having pseudoprincipal authors such as "anonFox."  $\text{op}_{d_2}$  only knows the state of posts after  $\text{op}_{d_1}$  occurs, and does not remove any posts because no post has author "Bea".

The choice of whether  $x_{[\mathbf{A}(d_1), \mathbf{A}(d_2)]}$  should result in  $\mathcal{S}_1$  or  $\mathcal{S}_2$  depends on the specific application and is left to the developer to specify.

In certain scenarios, however, Edna can only achieve end state  $\mathcal{S}_2$ . In order to achieve  $\mathcal{S}_1$ ,  $d_2$  must evaluate predicates against  $x_{\text{start}}$ , which requires knowing the value of  $x_{\text{start}}$ . To learn  $x_{\text{start}}$ , Edna must know how  $d_1$  modified  $x$  (e.g., knowing that a post with author "anonFox" originally was a post with author "Bea"); this information, however, is stored in the diffs produced by  $d_1$ .

Edna may not have access to these diffs when applying  $d_2$ : a client who invokes  $d_2$  while authenticated as principal  $q \neq p$ , and who does not provide the relevant capability pairs  $(C_{pd_1}^{\text{data}}, C_{pd_1}^{\text{loc}})$  cannot access  $p$ 's diffs for  $d_1$ .

In this case,  $d_2$  can only match against  $x_{[\mathbf{A}(d_1)]}$  and achieve end state  $\mathcal{S}_2$ . This is shown in Table 3.

	Access $d_1$ 's diffs while $\mathbf{A}(d_2)$ ?	
	Yes	No
$x_{[\mathbf{A}(d_1), \mathbf{A}(d_2)]}$	$\mathcal{S}_1$ or $\mathcal{S}_2$	$\mathcal{S}_2$

Table 3: Possibilities for  $x_{[\mathbf{A}(d_1), \mathbf{A}(d_2)]}$  depending on whether Edna has access to diffs from  $d_1$  regarding  $x$ .

**Semantics of Disguise Reversals.** We now consider composition of disguises in the presence of disguise reversals.

	Access to $d_i$ 's diffs while $\mathbf{A}(d_i)$ ?	
	Yes	No
$x_{[\mathbf{A}(d_i), \mathbf{R}(d_i)]}$	$x_{\text{start}}$	$x_{[\mathbf{A}(d_i)]}$

Table 4: Possibilities for  $x_{[\mathbf{A}(d_i), \mathbf{R}(d_i)]}$  depending on whether Edna has access to diffs from  $d_i$  regarding  $x$ .

We first consider  $x_{[\mathbf{A}(d_i), \mathbf{R}(d_i)]}$ ; Table 4 illustrates how this state depends the authorization a client invoking  $\mathbf{R}(d_i)$  provides to Edna to access  $d_i$ 's diffs for  $x$ . As expected, the updates applied by  $d_i$  to  $x$  cannot be revealed when the relevant diffs are inaccessible.

$d_1$ and $d_2$ 's Updates	Access to $d_1$ 's diffs while $\mathbf{R}(d_1)$ ?	
	Yes	No
Independent	$x_{[\mathbf{A}(d_2)]}$	$x_{[\mathbf{A}(d_1), \mathbf{A}(d_2)]}$
Conflict	$x_{[\mathbf{A}(d_1), \mathbf{A}(d_2)]}$	$x_{[\mathbf{A}(d_1), \mathbf{A}(d_2)]}$

Table 5:  $x_{[\mathbf{A}(d_1), \mathbf{A}(d_2), \mathbf{R}(d_1)]}$  depending on whether Edna has access to  $d_1$ 's diffs when revealing  $d_1$ , and whether updates from  $d_1$  and  $d_2$  to  $x$  modify the same data.

We next consider  $x_{[\mathbf{A}(d_1), \mathbf{A}(d_2), \mathbf{R}(d_1)]}$ : the end state of  $x$  when a disguise  $d_1$  is revealed after a subsequent disguise  $d_2$  has been applied. Table 5 illustrates how the result depends on Edna's access to  $d_1$  diffs, and whether  $d_1$  and  $d_2$  performed conflicting modifications to  $x$ 's attributes.

When the client invoking  $\mathbf{R}(d_1)$  provides no authorization to access to  $d_1$ 's diffs, then the updates applied by  $d_1$  to  $x$  cannot be revealed. However, even if the client grants access to  $d_1$ 's diffs, Edna will fail to reveal  $d_1$ 's updates to  $x$  if  $d_2$  and  $d_1$  both modify  $x$  in a conflicting manner:  $\mathbf{R}(d_1)$  must not restore data that was also disguised by  $d_2$ , and until  $d_2$  is revealed, invocations of  $\mathbf{R}(d_1)$  will not reveal  $x_{\text{start}}$ .

#### 4.1 Composition Implementation Techniques

Edna's algorithm for disguising and disguise reversal in most scenarios is straightforward.

To apply disguise  $d_i$ , Edna applies each  $\text{op}_{d_i}$  to all data objects satisfying  $\text{op}_{d_i}$ 's predicate, while also taking into account the information in accessible diffs (to e.g., predicate against undisguised versions of objects). Each  $\text{op}_{d_i}$  produces one or more diffs, which Edna stores either globally or encrypted by a data capability at  $C_{pi}^{\text{loc}}$ .

To reveal disguise  $d_i$ , Edna reveals the undisguised version of data stored in accessible diffs corresponding to  $d_i$  by updating the relevant objects  $x$  in the database.

However, Edna must be careful of two potential problems: (1) Edna cannot accidentally reveal data that must be disguised; and (2) Edna cannot let global diffs recording updates to data  $x$  leak information if  $x$  is subsequently (privately) disguised.

**Preventing Accidental Data Revelation.** Edna records the disguised state of data  $x$  in each  $\Delta_{pd}$  recording an update performed by  $d_i$  to  $x$ . If the state of  $x$  does not match the disguised state in  $d_i$ 's diff for  $x$ , then Edna knows the diff records a stale and overwritten update to  $x$ , and refuses to reveal the undisguised state of  $x$  stored in the diff.

**Preventing Global Diff Information Leakage.** A disguise may optionally store diffs in global storage, accessible to Edna without any client authorization. Such a global disguise provides no privacy guarantees against an adversary who compromises the application server, but does transform the application database so external users see disguised data.

If  $d_1$  is a global disguise and  $d_2$  a normal, privacy-preserving disguise,  $\text{op}_{d_2}$  may update a data object  $x$  previously updated by  $\text{op}_{d_1}$ .

$\text{op}_{d_1}$  produces a global diff  $\Delta_{pd_1}$ ;  $\text{op}_{d_2}$  produces a normal, encrypted and secured diff  $\Delta_{pd_2}$ .

If Edna simply performs  $\text{op}_{d_2}$ 's update to  $x$ , an adversary can still read undisguised data from the  $\Delta_{pd_1}$  diff because the diff is global and records a now-stale state of  $x$ ! To prevent global diffs like  $\Delta_{pd_1}$  from leaking data that should be disguised, Edna updates the data in  $\Delta_{pd_1}$  with  $\text{op}_{d_2}$ 's update. An update to a diff such as  $\Delta_{pd_1}$  itself generates a (private) diff  $\Delta_{pd_2}$ , just as updating a data object generates a diff.

Revealing  $d_2$  uses  $\Delta_{pd_2}$  to reveal the modification to  $\Delta_{pd_1}$ . If  $d_1$  has already been revealed, and the database state of  $x$  updated to reflect the contents of  $\Delta_{pd_1}$  (which were updated by the application of  $d_2$ ), then Edna simply uses  $\Delta_{pd_2}$  to reveal the current state of  $x$  in the database.

If a future  $d_3$  has further updated  $\Delta_{pd_1}$  in a way that conflicts with  $d_2$ 's update, then upon reversal of  $d_2$ , Edna will notice that  $\Delta_{pd_2}$  records an overwritten update to  $\Delta_{pd_1}$ , and will not revert the state of  $\Delta_{pd_1}$  until the future disguise  $d_3$  has been revealed.

[lyt: Note that removal is just an interesting case of this, in which  $\Delta_{pd_1}$  is removed, by a diff  $\Delta_{pd_2}$  stored in  $C_{pd_2}^{\text{loc}}$  that saves the removal of  $\Delta_{pd_1}$  for reversal.]